

Machine learning introduction

Part II – Deep Neural Networks

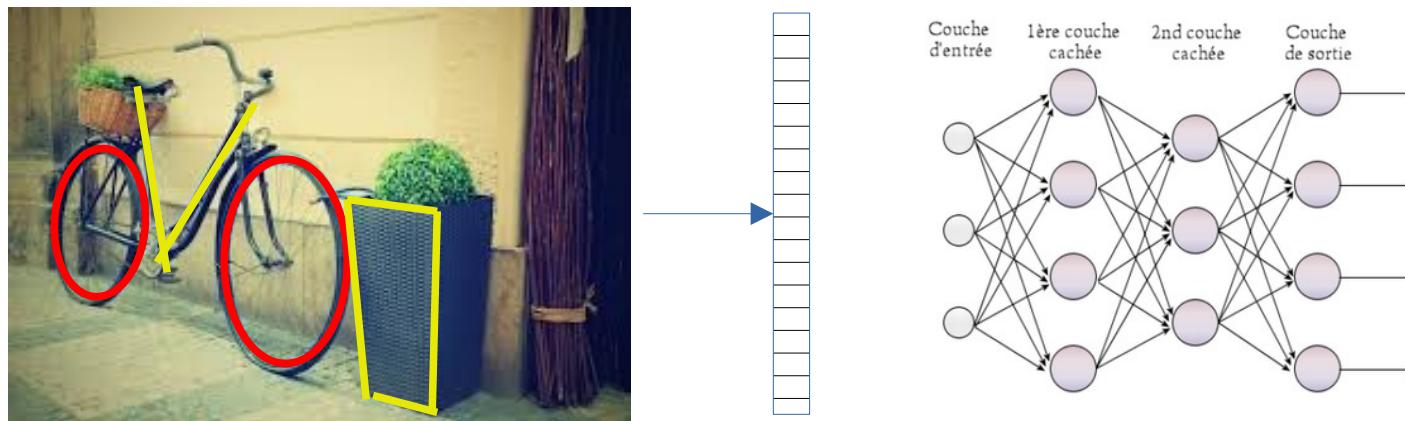
1 The deep networks

Simon Gay

Part II-1 : the deep networks

- The beginning: image recognition

- ‘Classic’ approach:
 - Detection of features (circles, lines, rectangles...)



- These features are used to fill a vector (number, position...)
 - A neural network is then trained on this vector
 - Not very reliable, error rate >25 % at ILSVRC 2011

Part II-1 : the deep networks

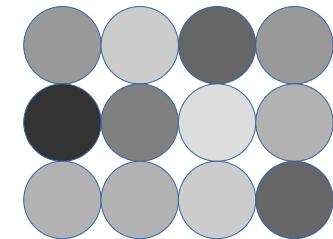
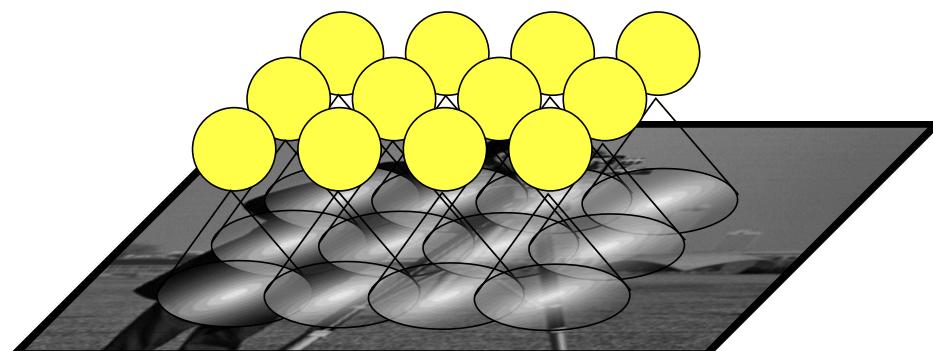
- **Deep learning**
 - Can a neural network generate such a feature vector from image ?
 - Problems:
 - More layers means more neurons and more connections to train
 - If the dataset is too small, the network would fail to generalize !
 - More neurons also means more computational resources
 - For many decades, neural networks were limited both by available resources and by available data sets.

Part II-1 : the deep networks

- **Deep learning**
 - Since middle of 2000s:
 - Computing power significantly increased
 - Parallelization on GPU (e.g. CUDA, 2007)
 - Dedicated architectures (Tensor Processing Unit, 2016)
 - But also collected data
 - GAFA and private life
 - Big data
 - Captcha
 - voluntarism
 - Open-source Dataset COCO (Common Objects in COntext) :
>200k labeled images, 80+ classes

Part II-1 : the deep networks

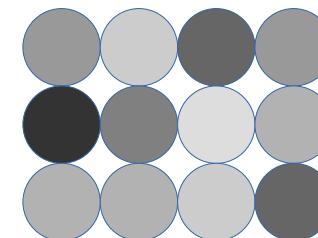
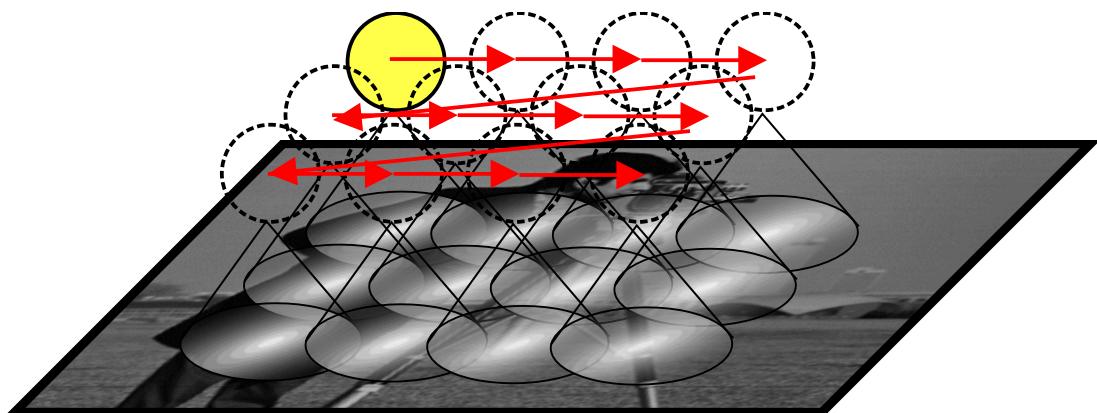
- Deep learning
 - But still not sufficient !
 - Example :
an image of 100x100 pixels requires 10 000 connections per input neuron !
 - Biological inspiration :
 - 1968: analysis on animals show that neurons of first processing layers are connected to small area of the retina
 - Several neurons detects the same feature



Detection of a feature
on the whole image

Part II-1 : the deep networks

- Deep learning
 - Yann Le Cunn (middle of 2000s) :
 - Idea: each neuron define a small feature and is applied on each position of the image



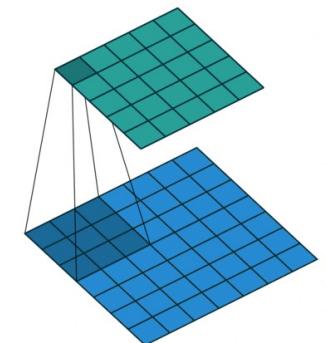
- Feature detection with few weights
- Works like a convolution → *convolutional neurons*

Part II-1 : the deep networks

- **Convolutions**

- A kernel

$$\begin{bmatrix} 0 & 0,25 & 0,5 & 0,25 & 0 \\ 0,25 & 0,5 & 0,75 & 0,5 & 0,25 \\ 0,5 & 0,75 & 1 & 0,75 & 0,5 \\ 0,25 & 0,5 & 0,75 & 0,5 & 0,25 \\ 0 & 0,25 & 0,5 & 0,25 & 0 \end{bmatrix}$$



<https://github.com/vdumoulin>

- Scalar product is applied on each position of the image
 - The result is a matrix giving the presence of kernel's feature on each position of the image



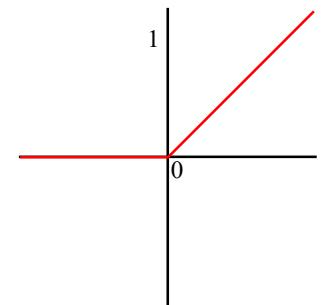
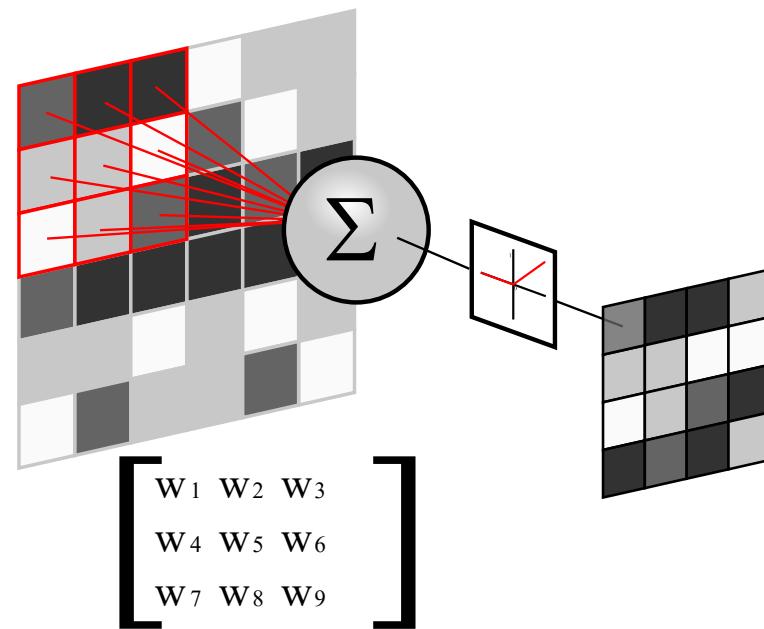
$$\otimes \begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix} =$$



Part II-1 : the deep networks

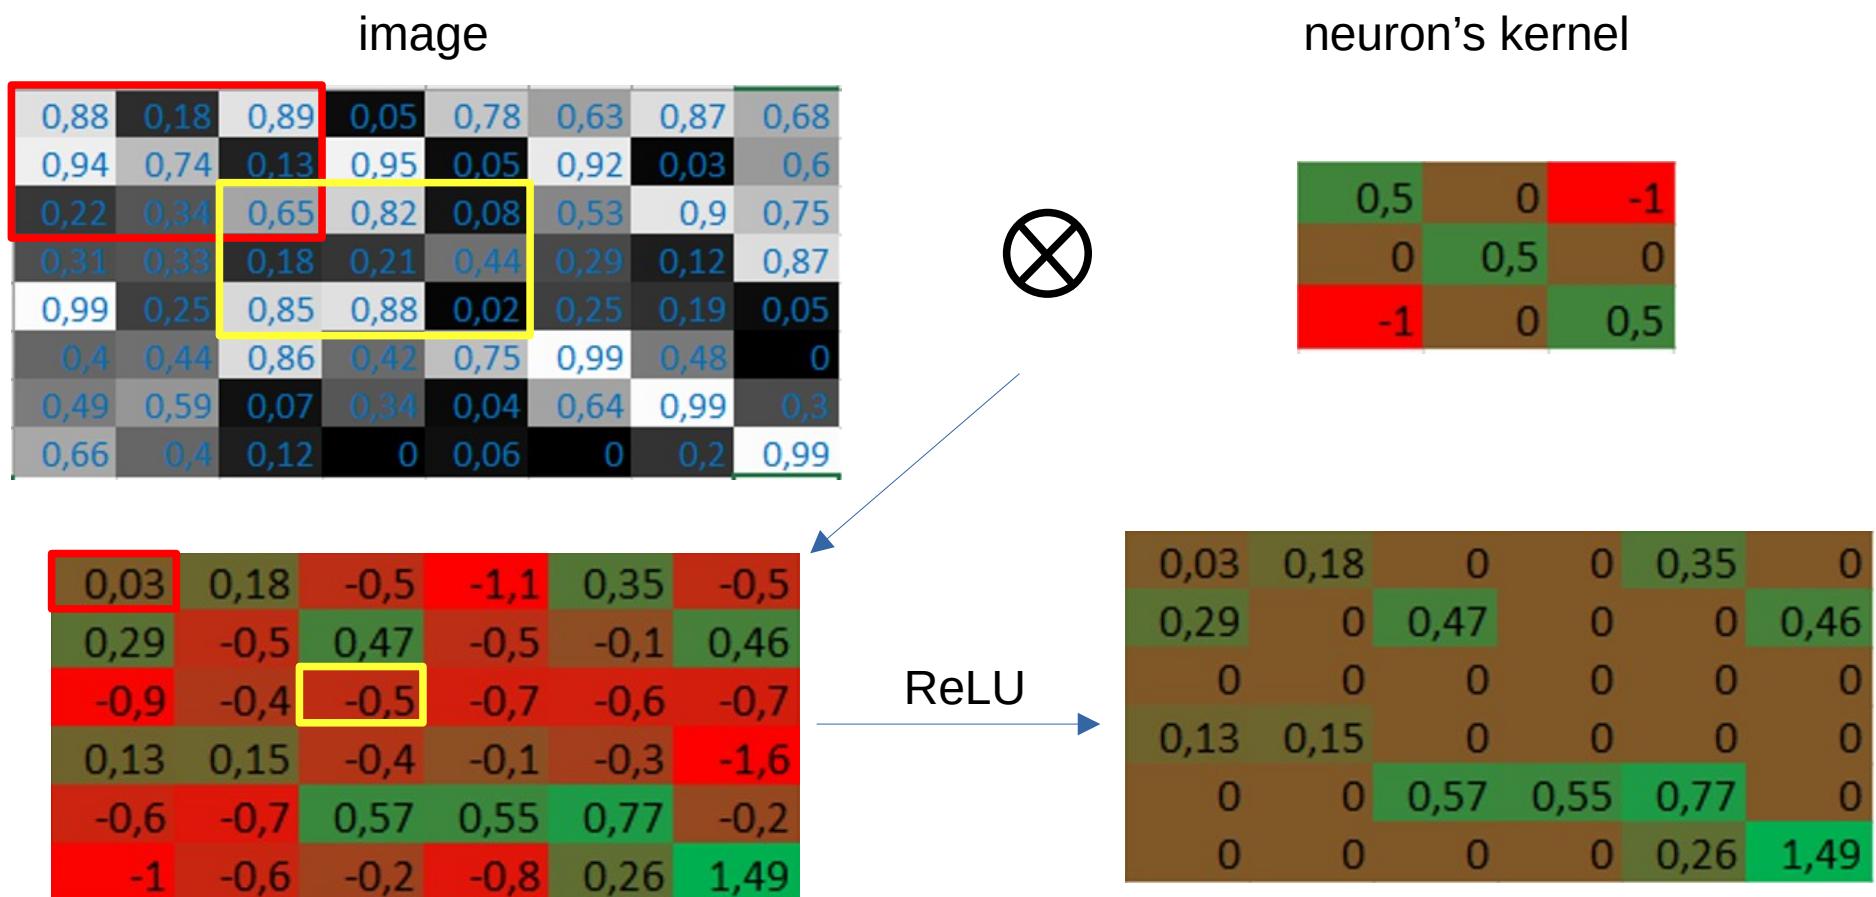
- **Convolutional neurons**

- Similar to a formal neuron
 - Input vector X
 - A set of weights W
→ kernel
 - An activation function
 - Reinforcement based on gradient descent
- Output is an “image”
 - Each pixel p_{xy} is the neuron output at position (x,y)
- Function RELU is often used : detect the presence of features



Part II-1 : the deep networks

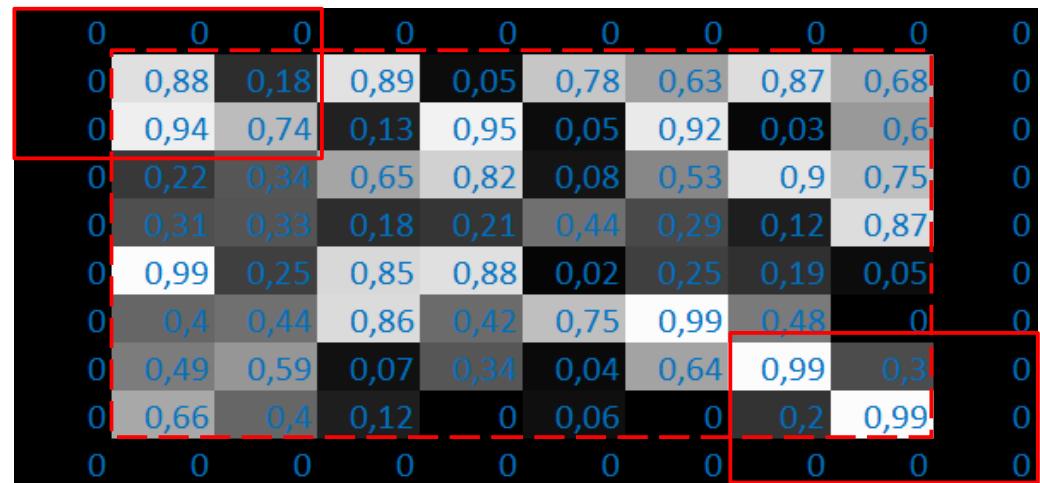
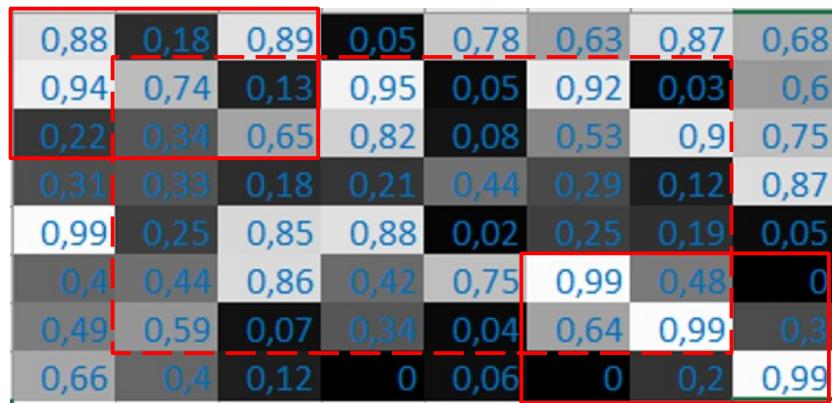
- **Convolutional neurons**



- We can notice that the output is smaller than the image

Part II-1 : the deep networks

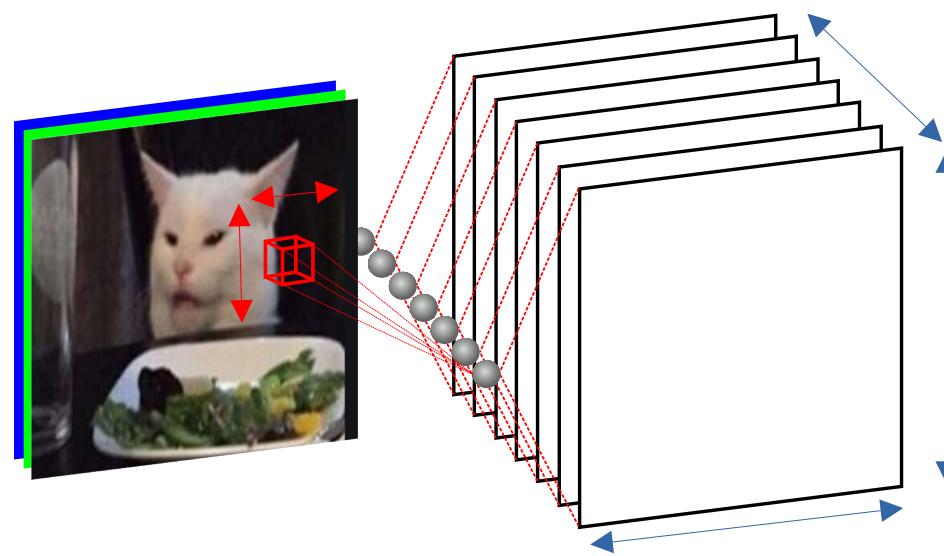
- **Deep Learning**
 - Padding :
 - Pixels with values of 0 are added around the image



Part II-1 : the deep networks

- **Deep learning**

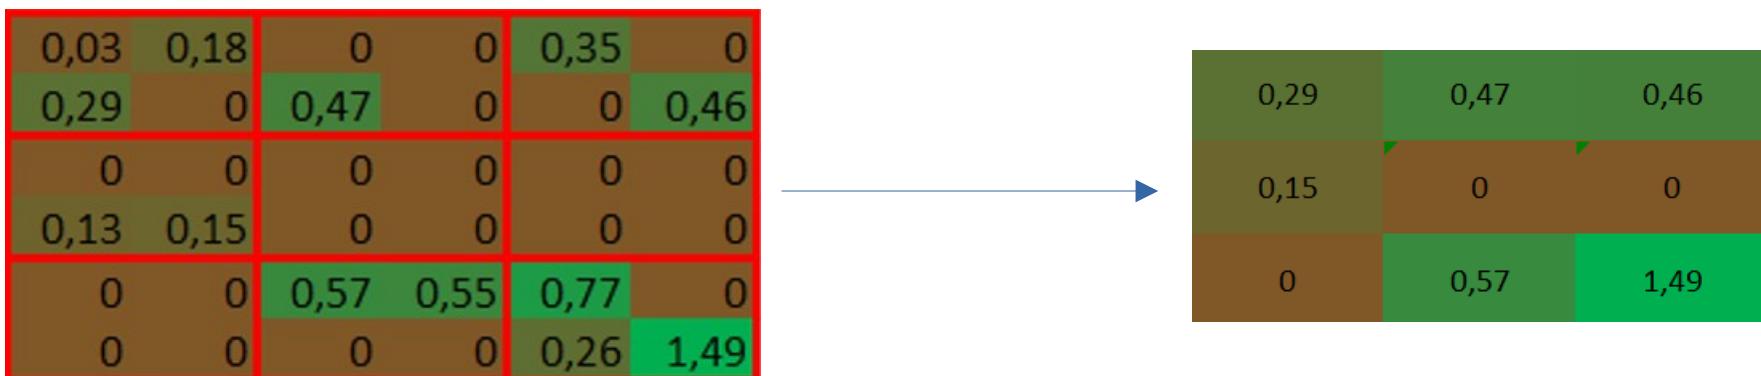
- Each neuron of the first layer generate its own convolution image
 - These images are gathered as a unique image with a depth
 - Depth of an image is the number of “layers”
 - A RGB image has a depth of 3
 - The kernel of a neuron has 3 dimensions (width, height and depth) and moves on 2 dimensions on the image



Part II-1 : the deep networks

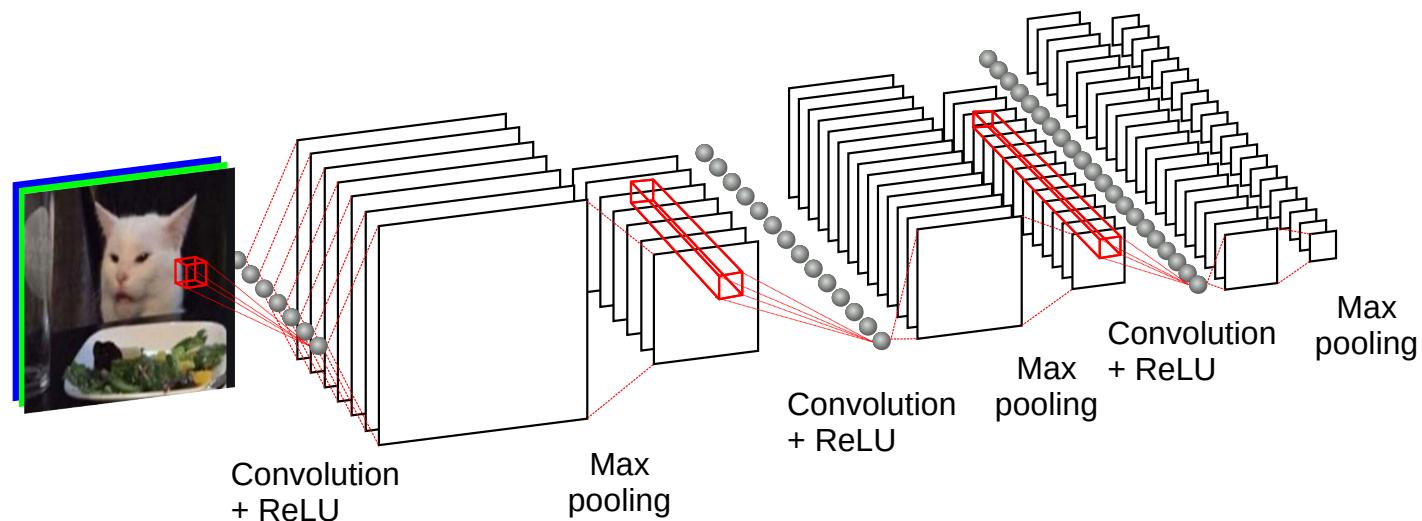
- **Pooling**

- As we looks for the presence of features, it is not necessary to keep adjacent pixels with high value
- Pooling layer :
 - Adjacent pixels are gathered
 - Average pooling → average of the group
 - Max pooling → keep maximum value (most used method)
 - Reduces the number of pixels for next layers



Part II-1 : the deep networks

- Deep learning
 - Convolution and pooling layers can be chained
 - We lose spatial information but get information about complex features

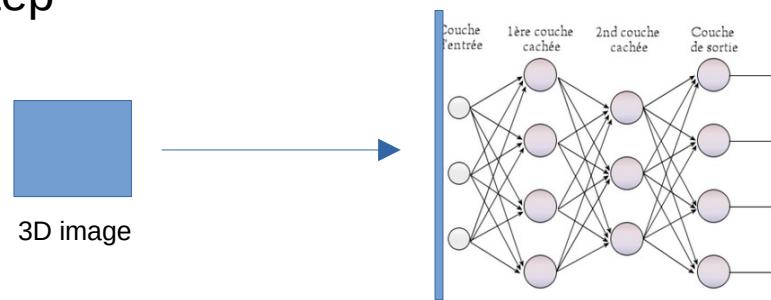


Part II-1 : the deep networks

- **Fully-connected layers**

- After several layers, the image looks more like a vector !
 - We thus can use it as input for a ‘normal’ neuronal network

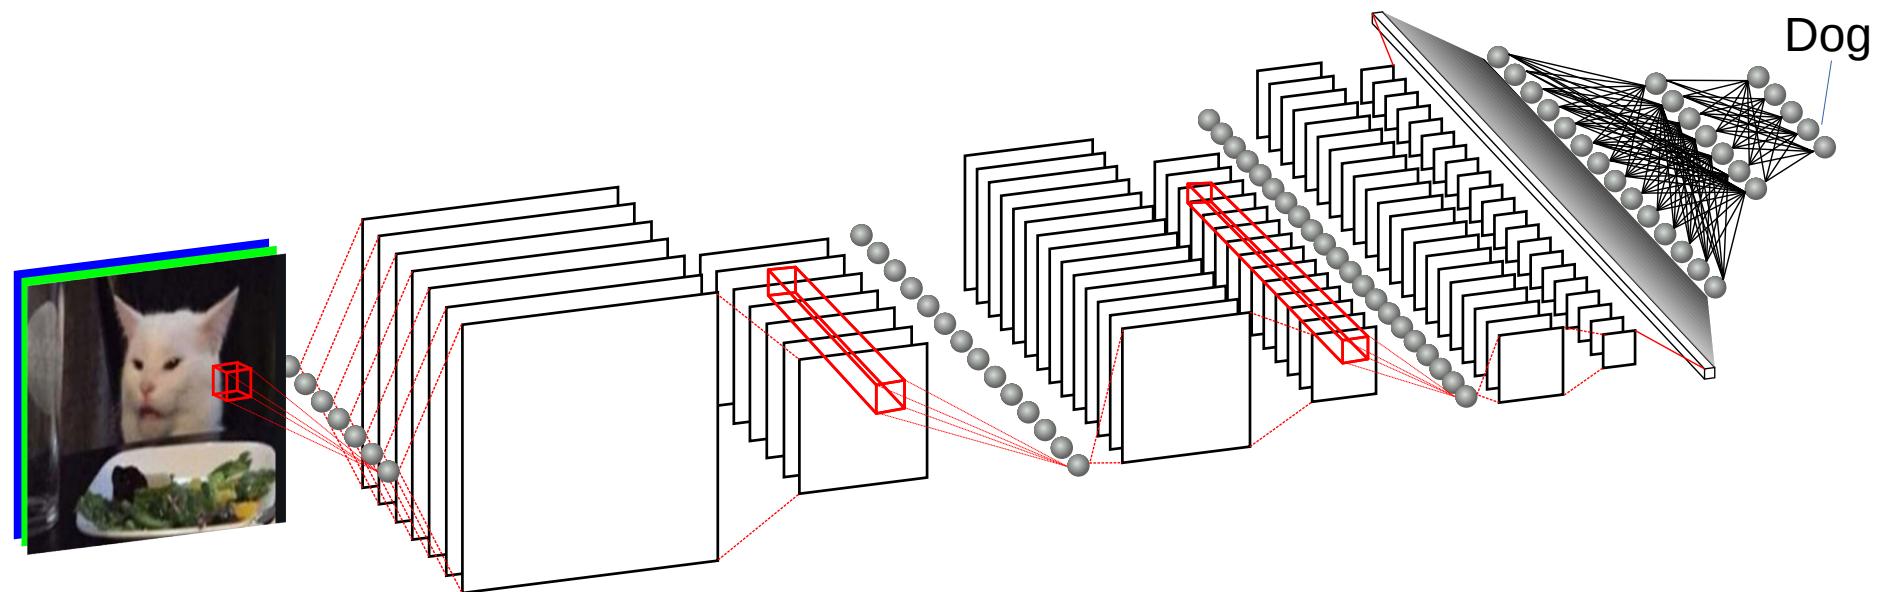
- Last image is converted to a 1-D vector
 - flattening step



- A neuronal network is then used to “interpret” the feature vector
 - these neuron layers are called ‘fully connected’

Part II-1 : the deep networks

- “Classical” architecture of a convolutional deep network

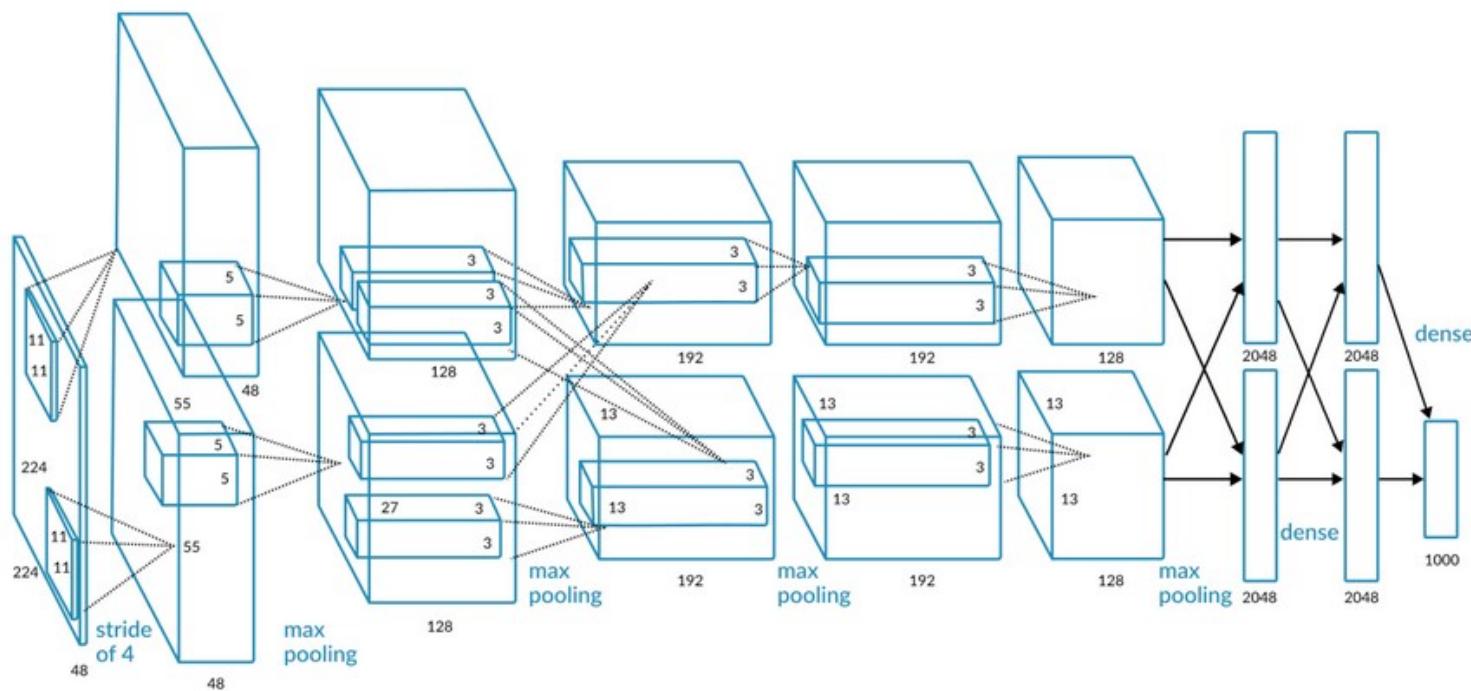


- convolution + ReLU
- Pooling
- convolution + ReLU
- Pooling
- ...
- Flattening
- fully connected layer
- Output layer

Part II-1 : the deep networks

- The architecture can be more complex:

example : AlphaNet (2012)



Part II-1 : the deep networks

- **Libraries for deep learning model development:**

- TensorFlow : developed by Google, open-source since 2015, derived from *Disbelief* project (2011)
 - Since 2017, a Lite version was created for embedded systems
- MXNet : Developed by Apache, can use many languages
- Caffe
- Theanos : since 2008
- Microsoft Cognitive Toolkit : since 2016
- PyTorch : developed by Facebook, based on Torch
- Keras : library of high level function to easily develop a deep network. Uses Tensorflow or Theanos libraries

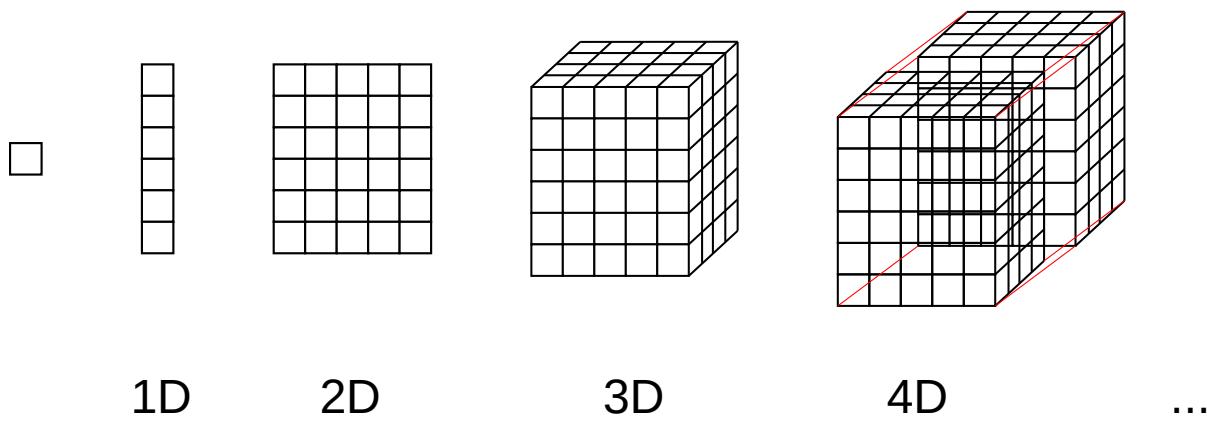


Part II-1 : the deep networks



- **TensorFlow**

- Library dedicated to the manipulation of *tensors*
 - A tensor is a matrix with a given number of dimensions



- Optimized functions
- Possibility to parallelize on GPU !

Part II-1 : the deep networks



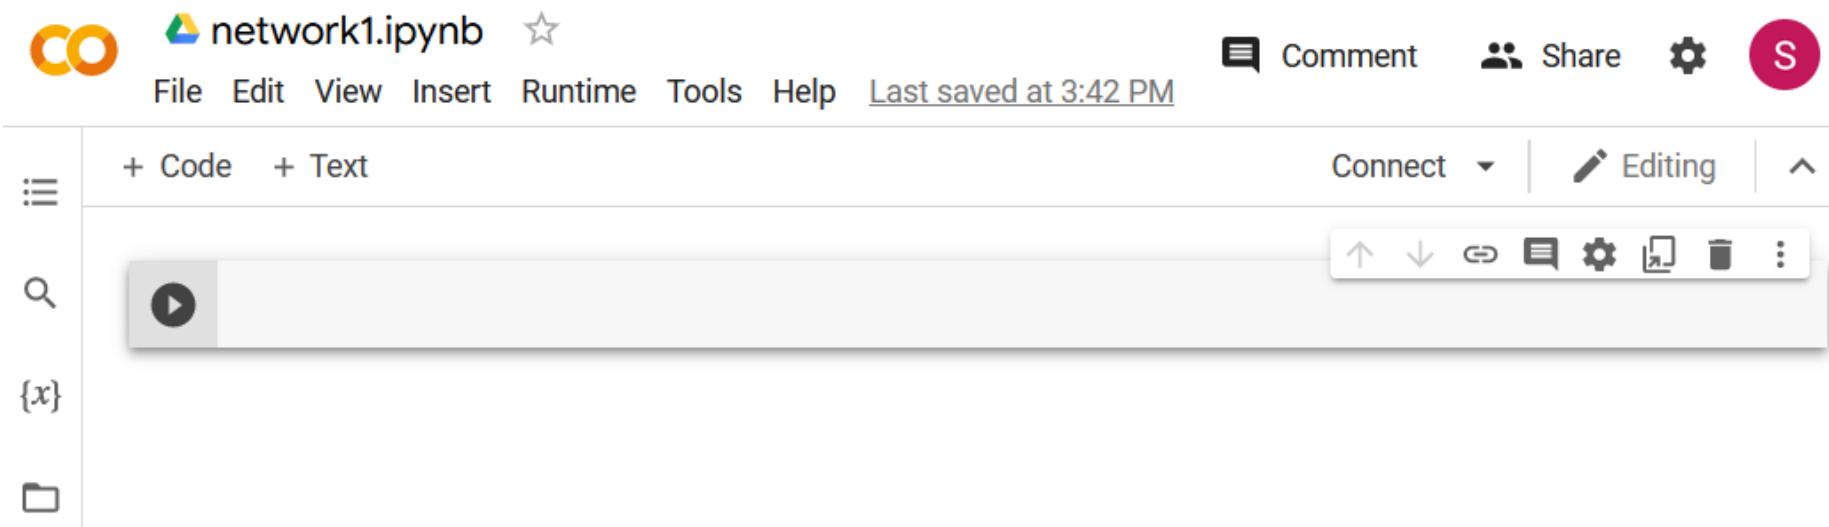
- **Keras**

- Library of high-level functions exploiting Tensorflow or Theanos
- Function to:
 - Load a ‘common’ dataset (ie used for benchmarking)
 - Define the architecture of a network
 - Define learning parameters and loss function
 - Train the network
 - Evaluate the network
 - Exploit the network on new data

Part II-1 : the deep networks



- **Let's start with Keras : implementing a convolutional network**
 - Open your session in Google Colab
 - Create a new notebook



Github : <https://github.com/gaysimon/ARTISAN2022>

Part II-1 : the deep networks



- **Let's start with Keras : implementing a convolutional network**
 - In the first cell, import the required libraries



```
import tensorflow as tf  
  
from tensorflow import keras  
  
import matplotlib.pyplot as plt  
import numpy as np
```

```
import tensorflow as tf  
  
from tensorflow import keras  
  
import matplotlib.pyplot as plt  
import numpy as np
```

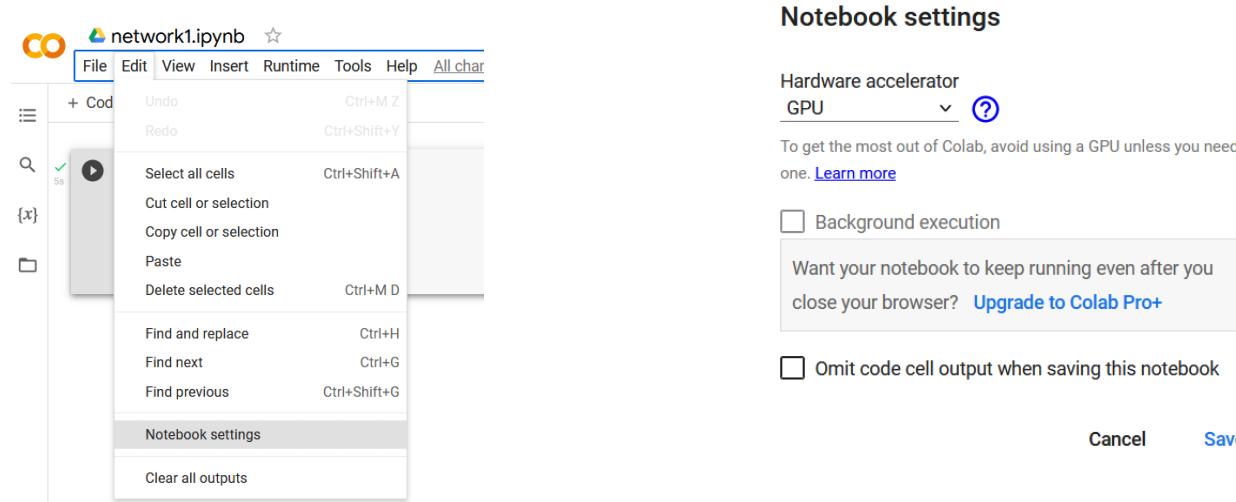
- Tensorflow and Keras for the neural network
- Pyplot and numpy for image display
- Run the cell to check the imports

Part II-1 : the deep networks



- **Using a GPU**

- Training a deep neural network requires a huge computational power, and can take several hours to train
- Hopefully, Tensorflow can parallelize calculations on a GPU
- Hopefully, Google Colab can provide GPUs
- Edit → notebook settings, then select ‘GPU’



Part II-1 : the deep networks



- **Using a GPU**

- DO NOT FORGET to turn off the GPU and stop the session after using it !
 - Runtime → manage sessions

Active sessions

Title	Last execution	RAM used	
 network1.ipynb Current session	GPU 0 minutes ago	0.89 GB	TERMINATE

- Otherwise, Google will reduce your access to GPUs for next sessions !

Part II-1 : the deep networks



- **Using a GPU**

- Then, copy this part of code and paste it just after importing Tensorflow

```
device_name = tf.test.gpu_device_name()
if device_name != '/device:GPU:0':
    raise SystemError('GPU device not found')

print('Found GPU at: {}'.format(device_name))
```

!\\ Python takes indentation into account

- If the message ‘found GPU’ appears, the GPU is ready

A screenshot of a Jupyter Notebook cell. The code imports tensorflow, checks for a GPU, and prints the GPU name if found. The output shows the message 'Found GPU at: /device:GPU:0'.

```
✓ 6s
import tensorflow as tf

device_name = tf.test.gpu_device_name()
if device_name != '/device:GPU:0':
    raise SystemError('GPU device not found')
print('Found GPU at: {}'.format(device_name))

from tensorflow import keras
import matplotlib.pyplot as plt
import numpy as np

Found GPU at: /device:GPU:0
```

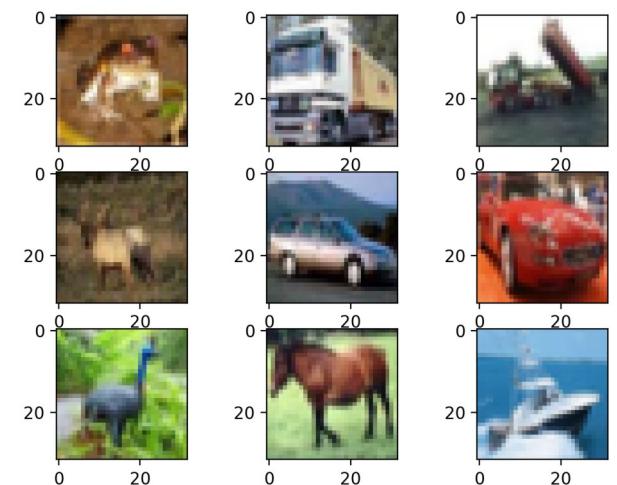
Part II-1 : the deep networks



- **Importing a dataset**

- Keras proposes different common datasets of different types :

- Boston housing prices (values)
→ *boston_housing*
 - CIFAR-10 (color images of size 32x32)
→ *cifar10*
 - CIFAR-100 (same but with 100 classes)
→ *cifar100*
 - Fashion MNIST (images B&W 28x28)
→ *fashion_mnist*
 - IMDB movie reviews (texts) → *imdb*
 - MNIST (images B&W 28x28) → *mnist*
 - Articles Reuters (texts) → *reuters*



CIFAR-10
(50 000 images)

Part II-1 : the deep networks



- **Importing a dataset**
 - We will use the MNIST dataset

- Add a new cell in your Colab ('+code')
- import dataset library
- Load the MNIST dataset into four matrices

```
▶ import tensorflow_datasets as tfds

dataset = keras.datasets.mnist
(img_train, label_train), (img_test, label_test) = dataset.load_data()
```

```
import tensorflow_datasets as tfds

dataset = keras.datasets.mnist
(img_train, label_train), (img_test, label_test) = dataset.load_data()
```

Part II-1 : the deep networks



- **Importing a dataset**

- The matrices must be converted
 - Images from integer [0,255] to float [0,1]
 - Labels must be converted from a number value to a vector of 10 results (e.g. 3 → [0, 0, 0, 1, 0, 0, 0, 0, 0, 0])
 - Keras proposes a function for that : *to_categorical*

```
import tensorflow_datasets as tfds
from tensorflow.keras.utils import to_categorical

dataset = keras.datasets.mnist
(img_train, label_train), (img_test, label_test) = dataset.load_data()

# convert images into float
img_train=img_train/255.0
img_test=img_test/255.0

output_train=keras.utils.to_categorical(label_train, num_classes=10)
output_test=keras.utils.to_categorical(label_test, num_classes=10)

print(img_train.shape)
print(img_test.shape)
```

(60000, 28, 28)
(10000, 28, 28)

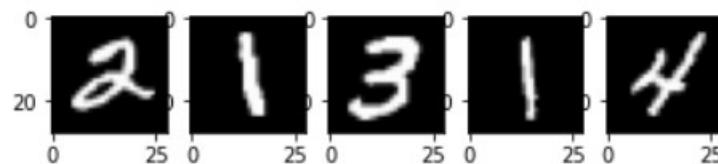
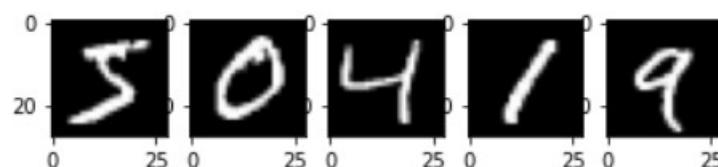
Part II-1 : the deep networks



- Importing a dataset
 - Let's display a sample of our dataset
 - Add a new colab cell
 - Copy the following code :

```
for i in range(10):
    plt.subplot(2,5,i+1)
    plt.imshow(img_train[i], cmap='gray')

plt.show()
```



Part II-1 : the deep networks



- **Importing a dataset**

- Problem : the images have 2 dimensions, but our network requires images with a depths !
- We add a dimension with *numpy's expand* function
 - Add a new colab cell
 - Add these lines :

```
img_train = np.expand_dims(img_train, axis=-1)
img_test = np.expand_dims(img_test, axis=-1)

print(img_train.shape)
print(img_test.shape)
```

(60000, 28, 28) → (60000, 28, 28, 1)
(10000, 28, 28) → (10000, 28, 28, 1)

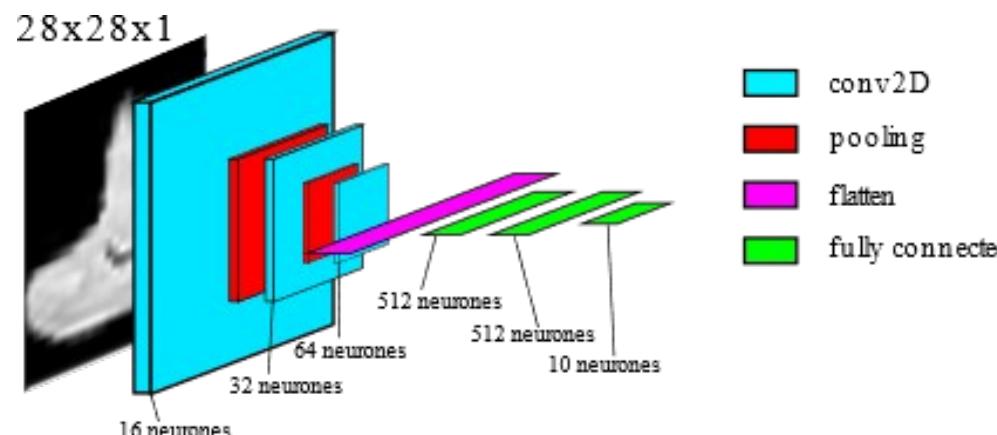
- Now, our images have a depth

Part II-1 : the deep networks



- **Implementing a convolutional network**

- We will implement this simple network :



- Create a new Colab cell
 - We start by declaring a Sequential network



```
model=keras.Sequential()
```

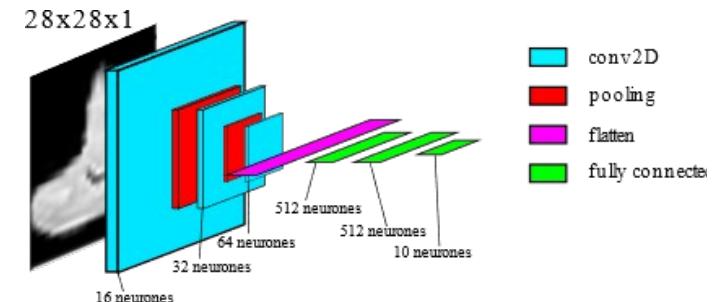
```
|
```

Part II-1 : the deep networks

• Implementing a convolutional network

- We then add layers :

- A first layer of convolutional neurons
 - Input images of size 28x28
 - 16 neurons
 - Convolution kernel of size 3x3
 - Padding
 - A ReLU activation function
- A max pooling layer with groups of size 2x2

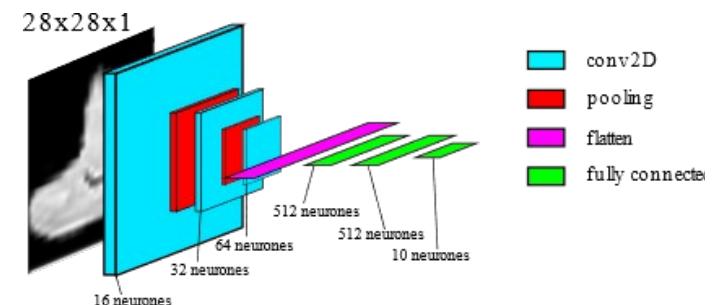


```
model=keras.Sequential()  
  
model.add(keras.layers.Conv2D(input_shape=(28,28,1), filters=16, kernel_size=(3,3), padding="same", activation="relu"))  
model.add(keras.layers.MaxPool2D(pool_size=(2,2), strides=(2,2)))  
|
```

```
model=keras.Sequential()  
  
model.add(keras.layers.Conv2D(input_shape=(28,28,1), filters=16, kernel_size=(3,3), padding="same", activation="relu"))  
model.add(keras.layers.MaxPool2D(pool_size=(2,2), strides=(2,2)))
```

Part II-1 : the deep networks

- **Implementing a convolutional network**
 - We then add layers :
 - A new convolutional layer (32 neurons)
 - A new max pooling
 - A last convolutional layer (64 neurons)

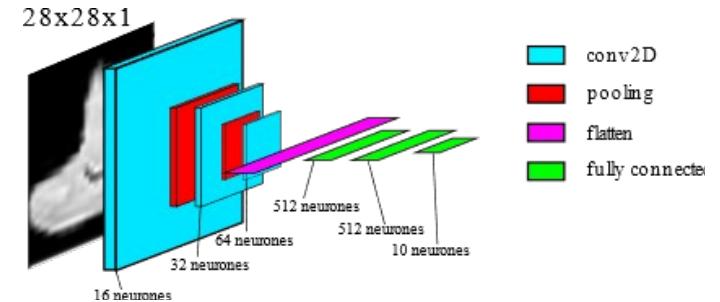


```
model=keras.Sequential()  
  
model.add(keras.layers.Conv2D(input_shape=(28,28,1), filters=16, kernel_size=(3,3), padding="same", activation="relu"))  
model.add(keras.layers.MaxPool2D(pool_size=(2,2), strides=(2,2)))  
  
model.add(keras.layers.Conv2D(filters=32, kernel_size=(3,3), padding="same", activation="relu"))  
model.add(keras.layers.MaxPool2D(pool_size=(2,2),strides=(2,2)))  
  
model.add(keras.layers.Conv2D(filters=64, kernel_size=(3,3), padding="same", activation="relu"))
```

Part II-1 : the deep networks

• Implementing a convolutional network

- We then add the fully connected part :
 - A flatten layer
 - 2 fully connected layers of 512 neurons
 - An output fully connected layer of 10 neurons (softmax function)



```
[ ]  
model=keras.Sequential()  
  
model.add(keras.layers.Conv2D(input_shape=(28,28,1), filters=16, kernel_size=(3,3), padding="same", activation="relu"))  
model.add(keras.layers.MaxPool2D(pool_size=(2,2), strides=(2,2)))  
model.add(keras.layers.Conv2D(filters=32, kernel_size=(3,3), padding="same", activation="relu"))  
model.add(keras.layers.MaxPool2D(pool_size=(2,2), strides=(2,2)))  
model.add(keras.layers.Conv2D(filters=64, kernel_size=(3,3), padding="same", activation="relu"))  
  
model.add(keras.layers.Flatten( ))  
  
model.add(keras.layers.Dense(units=512,activation="relu"))  
model.add(keras.layers.Dense(units=512,activation="relu"))  
  
model.add(keras.layers.Dense(units=10, activation="softmax"))
```

```
model.add(keras.layers.Flatten( ))  
  
model.add(keras.layers.Dense(units=512,activation="relu"))  
model.add(keras.layers.Dense(units=512,activation="relu"))  
  
model.add(keras.layers.Dense(units=10, activation="softmax"))
```

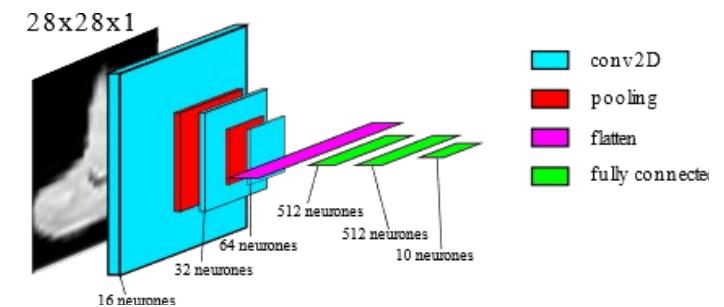
Part II-1 : the deep networks

- **Implementing a convolutional network**
 - The model architecture can be displayed with
model.summary()

```
model.summary()
```

```
→ Model: "sequential"
```

Layer (type)	Output Shape	Param #
<hr/>		
conv2d (Conv2D)	(None, 28, 28, 16)	160
max_pooling2d (MaxPooling2D)	(None, 14, 14, 16)	0
)		
conv2d_1 (Conv2D)	(None, 14, 14, 32)	4640
max_pooling2d_1 (MaxPooling2D)	(None, 7, 7, 32)	0
conv2d_2 (Conv2D)	(None, 7, 7, 64)	18496
flatten (Flatten)	(None, 3136)	0
dense (Dense)	(None, 512)	1606144
dense_1 (Dense)	(None, 512)	262656
dense_2 (Dense)	(None, 10)	5130
<hr/>		
Total params: 1,897,226		
Trainable params: 1,897,226		
Non-trainable params: 0		



Part II-1 : the deep networks



- **Keras** : some important types of layers (with some important parameters) :

- `tf.keras.layers.Dense(`

units,
activation=None,
use_bias=True,
)

→ number of neurons in the layer
→ activation function
→ use a bias or not

- `tf.keras.layers.Conv2D(`

filters,
kernel_size,
strides=(1, 1),
padding="valid",
activation=None,
use_bias=True,
)

note : versions 1D and 3D also exist

→ number of neurons
→ size of kernel
→ displacement (steps) of kernel
→ padding when "same"
→ activation function
→ use a bias or not

- `tf.keras.layers.MaxPooling2D(`

pool_size=(2, 2),
strides=None,
padding="valid"
)

→ pooling reduction
→ 'movements'

- Some activation functions: relu, sigmoid, softmax, tanh, ...
- Complete list of 100+ types of layers on: <https://keras.io/api/layers/>

Part II-1 : the deep networks



- **Keras** : construction of a network model
 - A sequential model can also be defined as a vector of layers:

```
model = Sequential([
    Conv2D(input_shape=(224,224,3), filters=64, kernel_size=(3,3), padding="same", activation="relu"),
    Conv2D(filters=64,kernel_size=(3,3), padding="same", activation="relu"),
    MaxPool2D(pool_size=(2,2), strides=(2,2)), ...
])
```

- Another method: functionnal model
 - Layers are created separately and connected

```
input_layer = Input(shape=(28,28,1))
conv1 = Conv2D(32, kernel_size=(3, 3), activation='relu', input_shape=(32,32,1))(input_layer)
conv2 = Conv2D(64, (3, 3), activation='relu', input_shape=(32,32,1))(input_layer)
merge = concatenate([conv1, conv2])
flatten = Flatten()(merge)
output_layer = Dense(num_classes, activation='softmax')(flatten)
model = Model(inputs=input_layer, outputs=output_layer)
```
 - Allows defining non sequential models (fusion, split...)

Part II-1 : the deep networks



- **Learning parameters**

- After constructing the network, we have to define
 - The loss function (error measure to optimize)
 - The optimization function (weight update algorithm)
- With a softmax activation function, we will use the CategoricalCrossentropy loss function !

Other loss functions :

- *MeanSquaredError* : $E = (r - y)^2$
- ... and dozens of other available functions : <https://keras.io/api/losses>
- In a new Colab cell, add :

A screenshot of a Jupyter Notebook cell. The cell contains a single line of Python code: "loss=keras.losses.CategoricalCrossentropy()". A play button icon is visible at the top left of the cell.

```
loss=keras.losses.CategoricalCrossentropy();
```

Part II-1 : the deep networks



- **Learning parameters**

- The optimization function is the function that reinforce neurons' weights
 - The SGD is the gradient descent presented previously
 - Other more optimized function adapt the learning rate to reduce learning time
 - Adam, Adadelta, Nadam, Ftrl... (see <https://keras.io/api/optimizers/>)
 - We select Adam optimizer, with a learning rate of 0.001 :



```
loss=keras.losses.CategoricalCrossentropy()  
optim=keras.optimizers.Adam(learning_rate=0.001)  
|
```

```
optim=keras.optimizers.Adam(learning_rate=0.001)
```

Part II-1 : the deep networks



- **Finalizing the network**

- Finally, the network is compiled : the neuron reinforcement functions are defined according to the network architecture and selected loss and optimization functions :



```
loss=keras.losses.CategoricalCrossentropy()  
optim=keras.optimizers.Adam(learning_rate=0.001)  
  
model.compile(loss=loss, optimizer=optim, metrics=["accuracy"])
```

```
model.compile(loss=loss, optimizer=optim, metrics=["accuracy"])
```

- Run the cell : if no error message appears, the network is ready !

Part II-1 : the deep networks



- **Training the network**

- Keras takes care of everything with function *fit*
 - Parameters:
 - The training data and label
 - The size of batches (number of data learned simultaneously)
 - The number of epoches
 - The level of details to display
 - In a new cell, write :



```
model.fit(img_train, output_train, batch_size=16, epochs=10, verbose=2)
```

```
Epoch 1/10  
3750/3750 - 13s - loss: 0.1175 - accuracy: 0.9638 - 13s/epoch - 3ms/step  
Epoch 2/10  
3750/3750 - 10s - loss: 0.0502 - accuracy: 0.9848 - 10s/epoch - 3ms/step  
-
```

- The learning process can takes several minutes...

Part II-1 : the deep networks



- **Evaluating/Exploiting the network**

- Keras takes care of everything (again) :
 - Function evaluate (works in a same way than fit)

```
▶ model.evaluate(img_test, output_test, batch_size=16, verbose=2)

625/625 - 1s - loss: 0.0703 - accuracy: 0.9884 - 1s/epoch - 2ms/step
[0.07027573138475418, 0.9883999824523926]
```

◀ < 1,2 %

- An image can be used as parameter to get a prediction :

```
▶ print(label_test[0:1])

model(img_test[0:1,:,:])
```

◀ [7] ←

```
<tf.Tensor: shape=(1, 10), dtype=float32, numpy=
array([[1.1164599e-26, 2.6863731e-18, 2.3253600e-17, 1.2962786e-17,
       2.5873384e-16, 5.7846084e-20, 1.6555734e-36, 1.0000000e+00,
       1.5967833e-22, 9.6893570e-12]], dtype=float32)>
```

Part II-1 : the deep networks

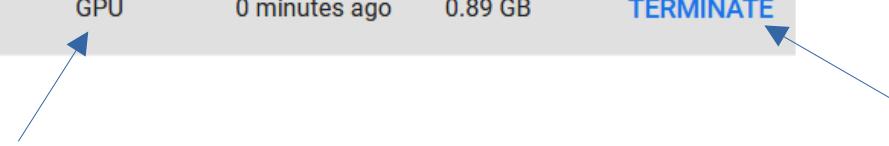


- Save your colab
- **Do not forget to disable the GPU and stop the session !**

Active sessions

Title	Last execution	RAM used
network1.ipynb Current session	0 minutes ago	0.89 GB

Edit → notebook settings, then select 'none'

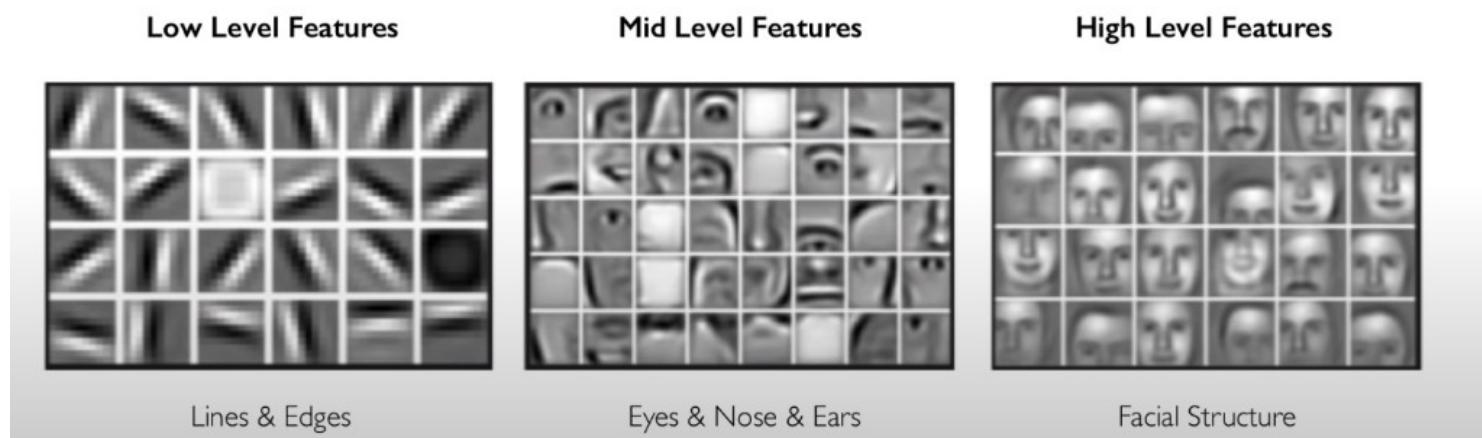


- **Note to go further :** there is a dataset similar to MNIST :
 - dataset = keras.datasets.**fashion_mnist**
 - You can try to improve your model to get the best accuracy

Training a deep network:
the question of the dataset

Part II-1 : the deep networks

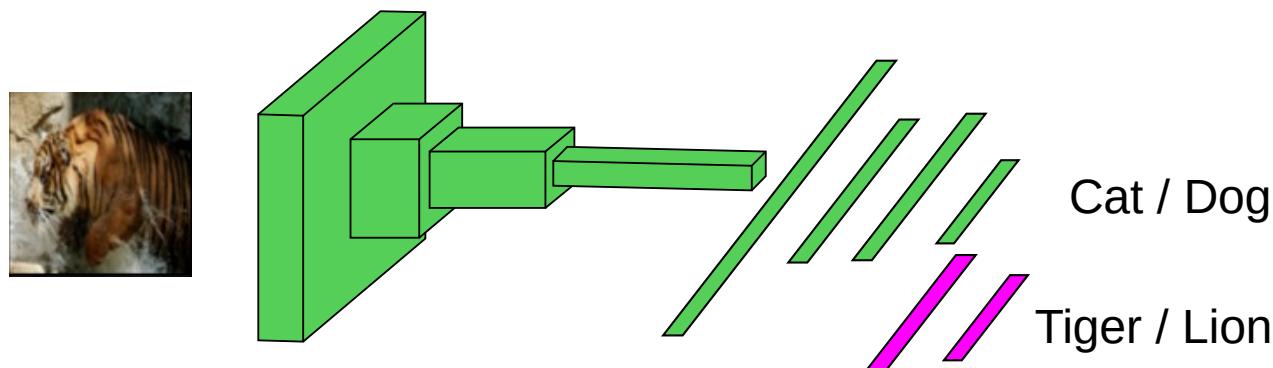
- **Training a deep network**
 - The more layers of neurons there are, the larger the learning dataset must be
 - A deep network may require a dataset of tens, even hundreds thousand images (or other kind of data)
 - Do we need a huge dataset for each new problem ?
 - For two similar problems, features defined by lowest layers are similar !



Part II-1 : the deep networks

- **Training a deep network**
 - For similar problem, it is possible to re-use low layers of an existing trained network
 - Principle of **transfer learning**

- We can use a network trained on a similar problem
- We then replace last layers of the network with new layers

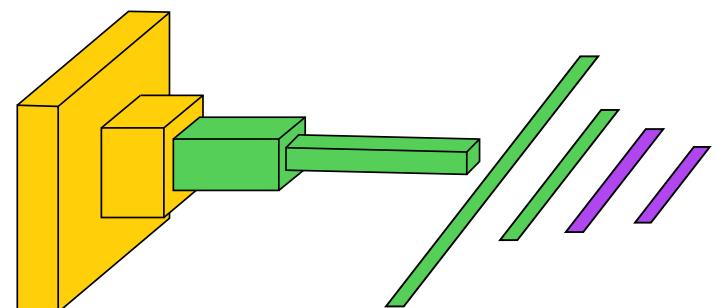
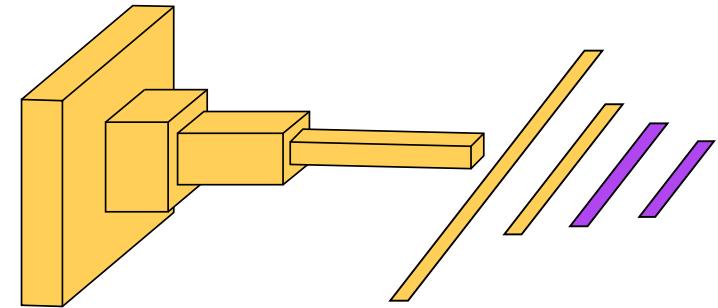
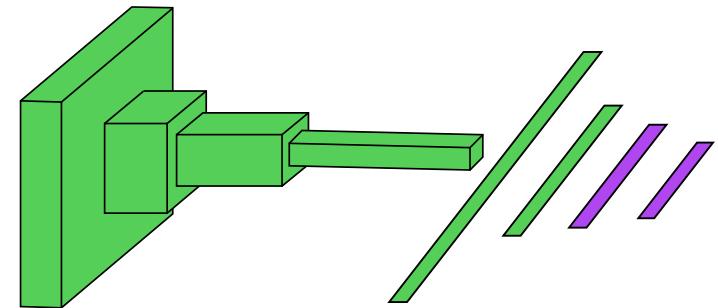


- The number of layers to replace depends on problems similarities and available dataset's size

Part II-1 : the deep networks

- **Transfer Learning**

- Fine tuning: last layers (often fully connected part) are replaced with new initialized layers
 - Faster learning with less examples
 - Low level features can adapt to the new problem
 - Still requires a large dataset
- Feature extraction: last layers are replaced and remaining layers are frozen
 - Fast learning with few examples
 - Resulting network is not always efficient
- Partial Fine tuning: olast layers are replaces, but only lowest layers are frozen
 - Good compromise between above methods
 - Low level features are more ‘universal’



Part II-1 : the deep networks

- **Transfer Learning**
 - **Several well-trained open-source networks are proposed for transfer learning :**
 - AlexNet : architecture used for ILSVRC 2012
 - VGGnet : one of most versatile. Uses 16 convolution layers
 - GoogleNet : winner of ILSVRC 2014, uses 22 convolution layers
 - ResNet : winner of ILSVRC 2015 (with an error ratio of 3,57 %, it is more efficient than human participants!), uses 152 convolution layers

Part II-1 : the deep networks



- **Let's implement a true deep neural network !**

- New way to build a dataset
- Transfer learning from an existing network
- Goal : recognize a cat from a dog
- Create a new Colab Notebook, enable GPU, and copy the imports

CO deep1.ipynb

File Edit View Insert Runtime Tools Help

+ Code + Text

```
import tensorflow as tf

device_name = tf.test.gpu_device_name()
if device_name != '/device:GPU:0':
    raise SystemError('GPU device not found')
print('Found GPU at: {}'.format(device_name))

from tensorflow import keras
import matplotlib.pyplot as plt
import numpy as np
```

Found GPU at: /device:GPU:0

```
import tensorflow as tf

device_name = tf.test.gpu_device_name()
if device_name != '/device:GPU:0':
    raise SystemError('GPU device not found')
print('Found GPU at: {}'.format(device_name))

from tensorflow import keras
import matplotlib.pyplot as plt
import numpy as np
```

Part II-1 : the deep networks



- **The dataset**

- We will download a set of images and unzip them in different folders
- Copy and paste these code lines in two different colab cells, and execute them :

```
!wget --no-check-certificate \
    https://storage.googleapis.com/mledu-datasets/cats_and_dogs_filtered.zip \
    -O /tmp/cats_and_dogs_filtered.zip
```

```
import os
import zipfile

local_zip = '/tmp/cats_and_dogs_filtered.zip'
zip_ref = zipfile.ZipFile(local_zip, 'r')
zip_ref.extractall('/tmp')
zip_ref.close()

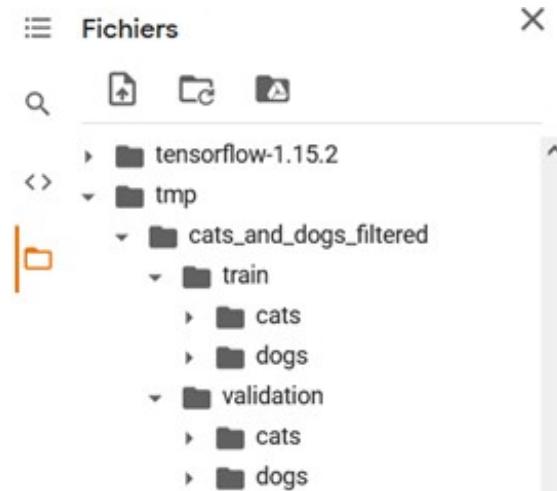
train_dir = '/tmp/cats_and_dogs_filtered/train'
test_dir = '/tmp/cats_and_dogs_filtered/validation'
```

Part II-1 : the deep networks



- **The dataset**

- The images are stored in a specific folder tree



- Keras proposes an object that can browse this kind of folder tree to feed a neural network for training or evaluating it :
 - The *ImageDataGenerator*

Part II-1 : the deep networks



- **The dataset**

- In a new Colab cell, we write the image generator :

```
trainDataGenerator = keras.preprocessing.image.ImageDataGenerator(rescale=1./255)
trainData = trainDataGenerator.flow_from_directory(directory=train_dir, target_size=(224,224))

testDataGenerator = keras.preprocessing.image.ImageDataGenerator(rescale=1./255)
 testData = testDataGenerator.flow_from_directory(directory=test_dir, target_size=(224,224))
```

Found 2000 images belonging to 2 classes.

Found 1000 images belonging to 2 classes.

- Then, we load the VGG16 network

```
VGG16= keras.applications.VGG16(weights="imagenet")
VGG16.summary()
```

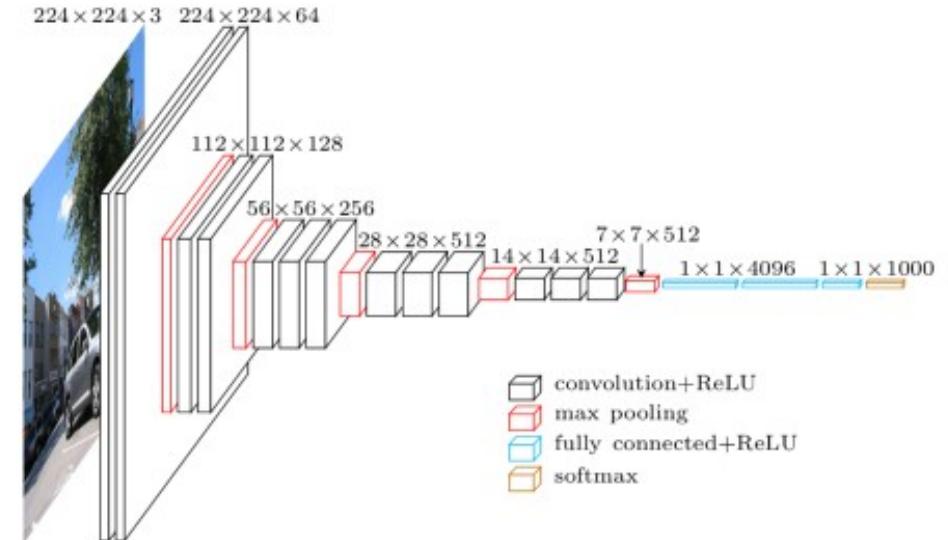
Part II-1 : the deep networks



- **The VGG16 network**

- We cannot train a network with so many parameters

Total params: 138,357,544
Trainable params: 138,357,544
Non-trainable params: 0



- Especially with such a small dataset !
 - Hopefully, the VGG16 network was trained with multiple animals (including cats and dogs) : the lowest layers may recognize pertinent features
- Note : VGG16 was trained during multiple days on high-end GPUs with *imageNet* dataset (>1M labeled images)

Part II-1 : the deep networks



- **The VGG16 network**

- The VGG16 network, as proposed by Keras, can be loaded without the fully connected part
 - Parameter `include_top` can be set to `False`
 - Parameter `weights='imagenet'` allows loading weight obtained with imagenet dataset (10M+ images)
- Add a new Colab cell and paste this :

```
base_VGG16= keras.applications.VGG16(weights="imagenet", include_top=False, input_shape=(224,224,3))
base_VGG16.summary()
```



- Observe where the network is cut :
 - As the flatten layer is not included, other convolutional layers can be added

Part II-1 : the deep networks



- **The VGG16 network**

- We freeze the weights of the network (in a new Colab cell) :

```
base_VGG16.trainable=False  
  
base_VGG16.summary()
```

```
=====  
Total params: 14,714,688  
Trainable params: 0 ←  
Non-trainable params: 14,714,688
```

Part II-1 : the deep networks



- **The VGG16 network**

- The VGG16 part is then used as a layer for our network
- We then add the fully connected part using a flatten layer, a fully connected layer with 50 neurons and an output flatten layer with 2 neurons

```
model=keras.models.Sequential()
model.add(base_VGG16)

model.add(keras.layers.Flatten())
model.add(keras.layers.Dense(50, activation="relu"))
model.add(keras.layers.Dense(2, activation="softmax"))

model.summary()
```

```
=====
Total params: 15,969,240
Trainable params: 1,254,552 ←
Non-trainable params: 14,714,688
```

Part II-1 : the deep networks



- **The VGG16 network**

- The network is compiled with Categorical_Crossentropy and Adam functions

```
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

- And trained using the data generator that provides both data and labels

```
model.fit(trainData, epochs=5, batch_size=32)
```



```
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

model.fit(trainData, epochs=5, batch_size=32)
```

Epoch 1/5

19/63 [=====>.....] - ETA: 6s - loss: 0.8847 - accuracy: 0.7105

Part II-1 : the deep networks



- **The VGG16 network**

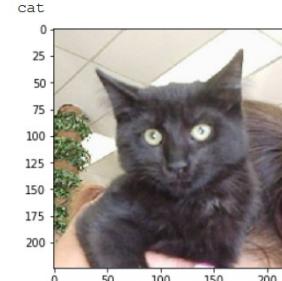
- The network can then be evaluated :

```
model.evaluate(testData, batch_size=32, verbose=2)
```

```
32/32 - 6s - loss: 0.2513 - accuracy: 0.9080 - 6s/epoch - 186ms/step  
[0.2513466775417328, 0.9079999923706055]
```

- The following code allows testing an image :

```
test_cats_files = os.listdir(test_dir+"/cats")  
test_dogs_files = os.listdir(test_dir+"/dogs")  
  
img = keras.preprocessing.image.load_img(test_dir+"/cats/"+test_cats_files[8],target_size=(224,224))  
#img = keras.preprocessing.image.load_img(test_dir+"/dogs/"+test_dogs_files[8],target_size=(224,224))  
img = np.asarray(img)  
plt.imshow(img)  
img = np.expand_dims(img, axis=0)  
  
prediction=model.predict(img)  
  
if prediction.argmax()==0: print("cat")  
else : print("dog")
```



Part II-1 : the deep networks

- **Applications**
 - Question : is it possible to recreate the image from the feature vector ?
 - Successive convolution generate an information loss
 - The feature vector only contains a fraction of image information
 - New layers to reverse convolutional layer's processing
 - De-pooling
 - transposed convolution

Part II-1 : the deep networks

- **Image reconstruction**

- De-Pooling :

- Nearest neighbor

0,29	0,47	0,46
0,15	0	0
0	0,57	1,49

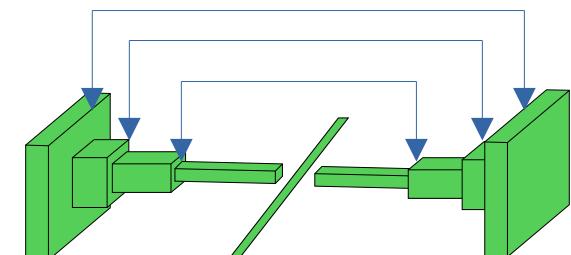
0,29	0,29	0,47	0,47	0,46	0,46
0,29	0,29	0,47	0,47	0,46	0,46
0,15	0,15	0	0	0	0
0,15	0,15	0	0	0	0
0	0	0,57	0,57	1,49	1,49
0	0	0,57	0,57	1,49	1,49

- Bed of nails

0,29	0	0,47	0	0,46	0
0	0	0	0	0	0
0,15	0	0	0	0	0
0	0	0	0	0	0
0	0	0,57	0	1,49	0
0	0	0	0	0	0

- Max UnPooling

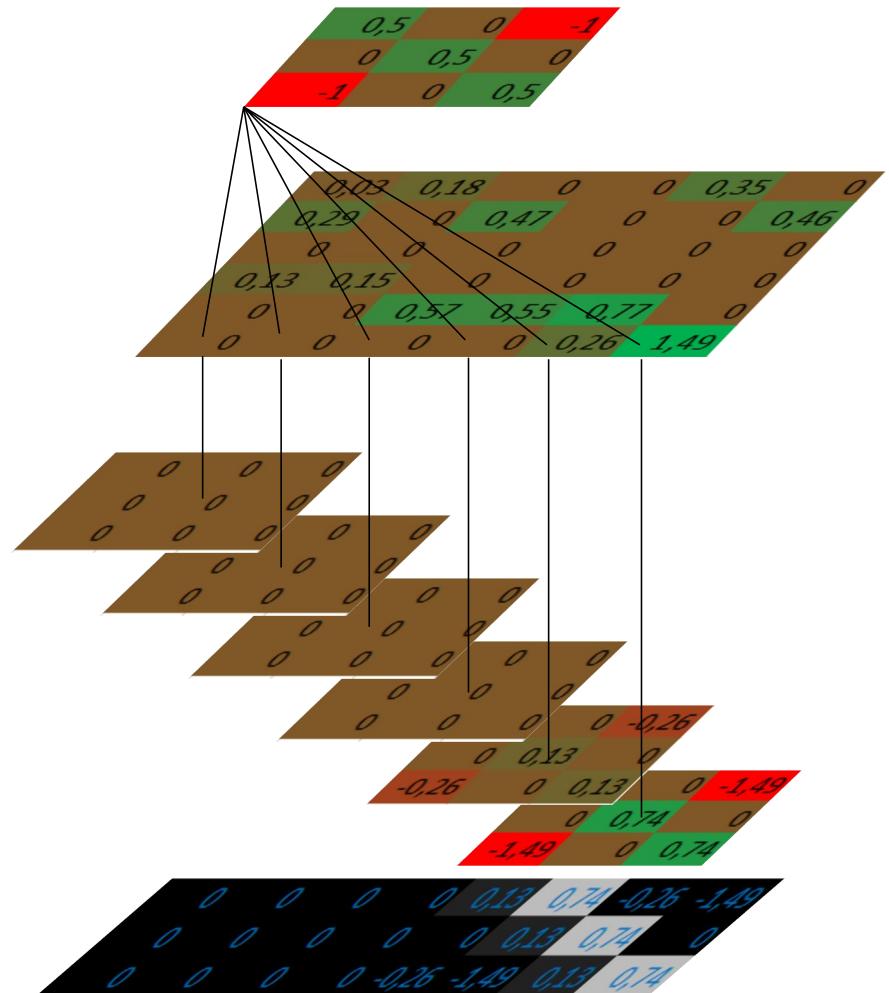
- use the mirror pooling layer to define which pixel get the value



Part II-1 : the deep networks

- **Image Reconstruction**

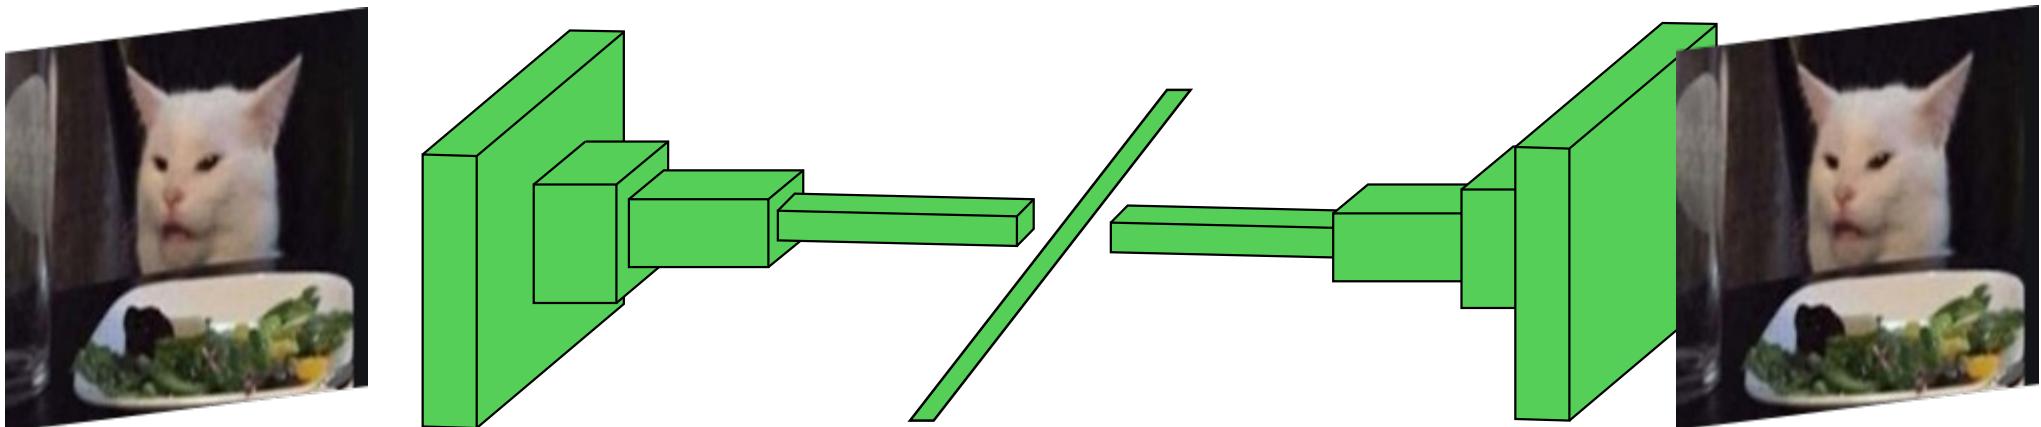
- Transposed Convolution:
 - Principle of convolution is ‘reversed’
 - Output image is the sum of kernel weighted by convolution image values



Part II-1 : the deep networks

- **Image reconstruction**

- The fully-connected part is replaced by a ‘mirror’ version of the network

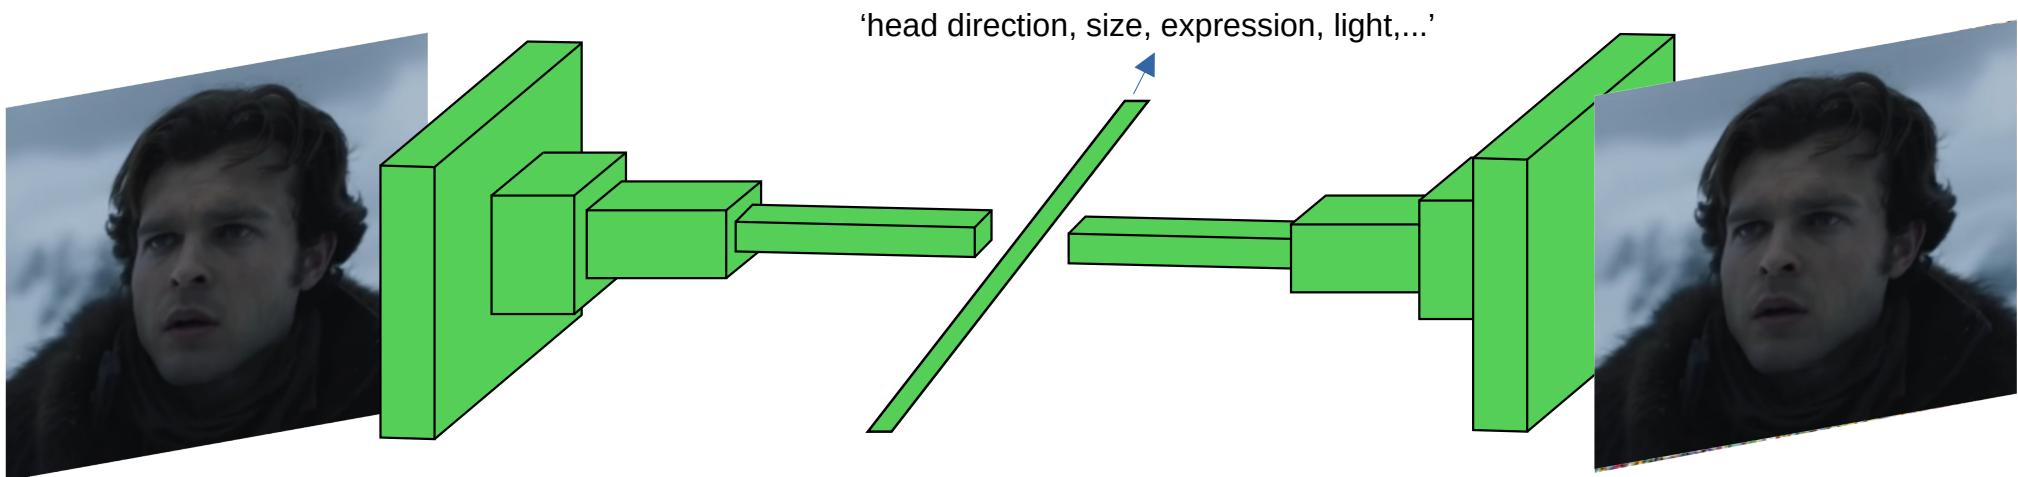


- Input image is used as expected result ($E = | \text{Img}_{\text{ref}} - \text{Img}_{\text{gen}} |^2$)
 - First part encodes the information: ‘encoder’ network
 - Second part ‘decode’ the feature vector to reconstruct the image: ‘decoder’ network
 - Decoder learns to generate textures, edges... from feature description
 - Very efficient way to de-noise an image

Part II-1 : the deep networks

- **Image reconstruction**

- The network ‘learns’ to define minimal information
- Let’s train the network on images of a specific face

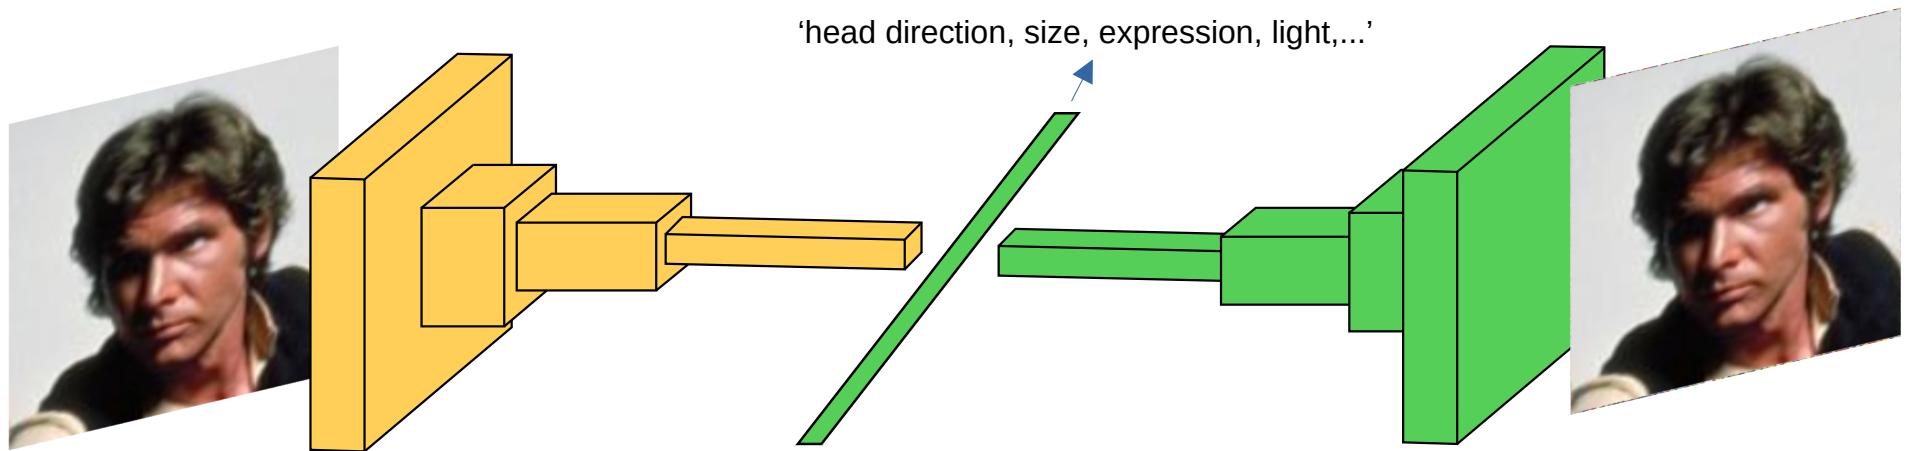


- The network will learn to encode minimum data to reconstruct the face
 - Head orientation and size, facial expression... (implicitly)
 - The decoder learns to reconstruct the image with these features

Part II-1 : the deep networks

- **Image reconstruction**

- The encoder's weights are frozen
- Decoder is reinitialized and trained with another face

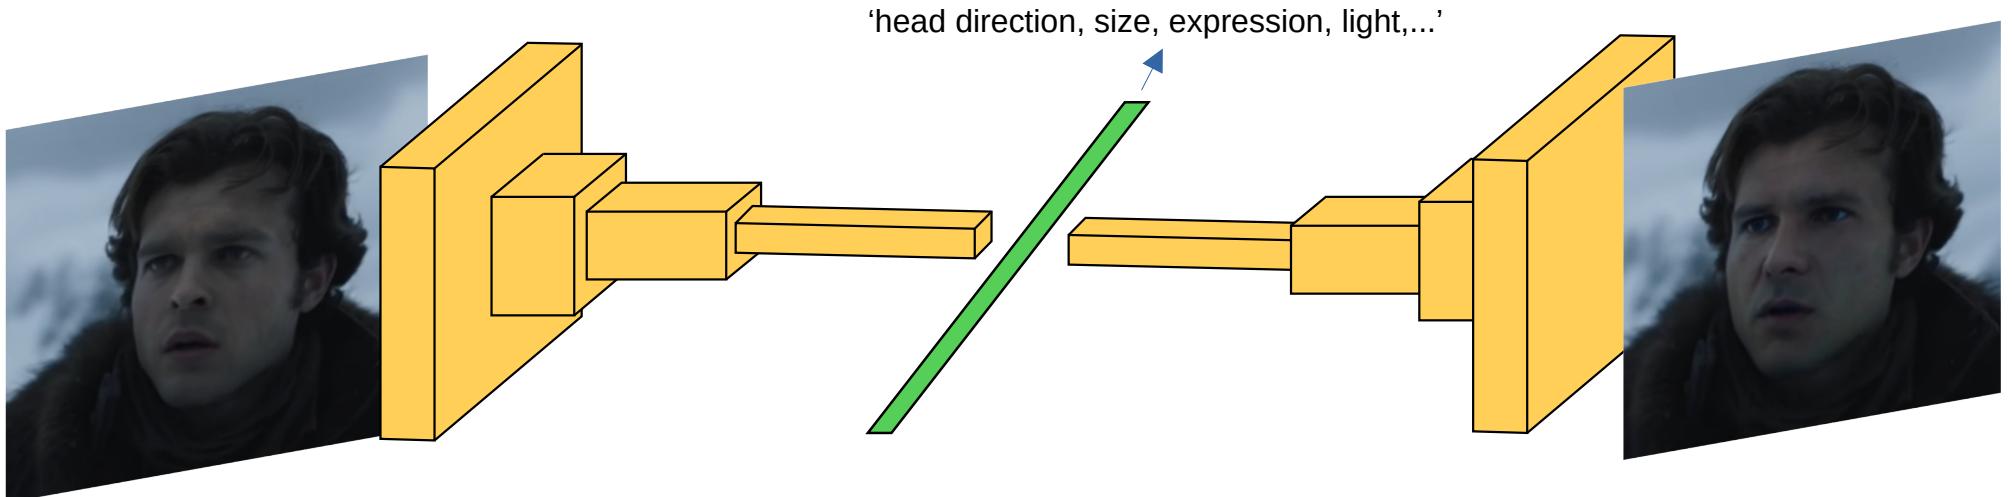


- The encoder will provide feature vectors in a similar way
- But the decoder must learn to generate this face from features

Part II-1 : the deep networks

- **Image reconstruction**

- We now froze the whole network
- An image of first face is presented to the network
 - The encoder extracts features from this face
 - The decoder will generate an image of the second face from these features



- Principle of deep fake

Part II-1 : the deep networks

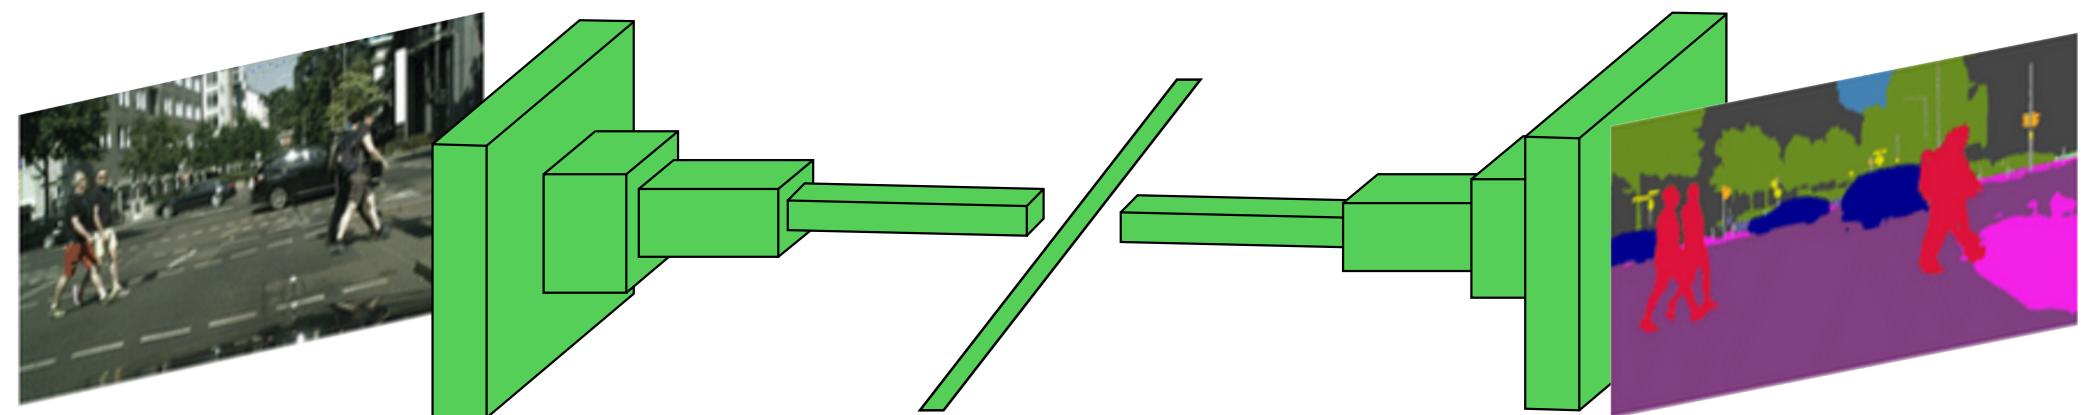
- Image reconstruction



Harrison Ford in Solo: A Star Wars Story (Shamook)

Part II-1 : the deep networks

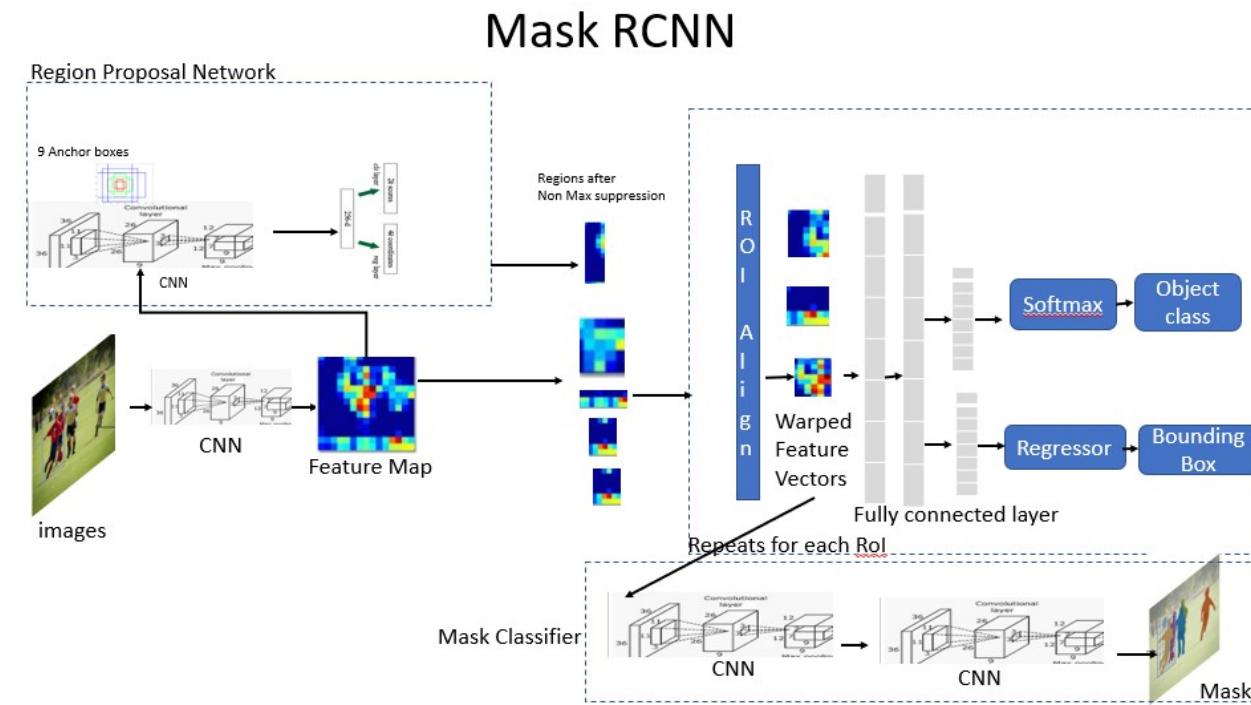
- **Fully-convolutional networks**
 - We can also use other type of images as output
 - Example : Mask-RCNN
 - The network is trained on images, with hand-labeled mask images



Part II-1 : the deep networks

- **Fully-convolutional network**

- Mask-RCNN : in practice, a little more complex
 - Detection of objects in the image
 - extraction of region of interest containing objects
 - Detection pixel by pixel to generate a mask



Part II-1 : the deep networks

- **Fully-convolutional network**
 - Mask-RCNN, trained on dataset COCO (Common Objects in COnext)



Training a deep network: bias and security

Part II-1 : the deep networks

- **The dataset bias**
 - Training a deep network requires a huge number of examples
 - Gathering a large number of example can be difficult, leading to statistical bias in the collected samples
 - Neural networks are very sensitive to statistical bias in datasets and will use ANY regularity they can found to minimize errors
 - Example : a network defining the age of a person uses the size of ears
 - These bias may generate serious issues
 - In 2014, Amazon used a deep network to analyze applicants' resume
 - This network massively rejected women's resume
 - The network was trained on resumes collected during 10 past years, most of them were from men...

Part II-1 : the deep networks

- **The dataset bias**

- Example de biais :
 - Network trained to recognize wolves from husky
 - High success ratio
 - But some (obvious) errors !
- After analyzing features, it appeared that the network only detect snow on the background !
 - In the dataset, a large majority of wolf pictures have snow in the background



Predicted: **wolf**
True: **wolf**



Predicted: **husky**
True: **husky**



Predicted: **wolf**
True: **wolf**



Predicted: **husky**
True: **husky**



Predicted: **wolf**
True: **wolf**

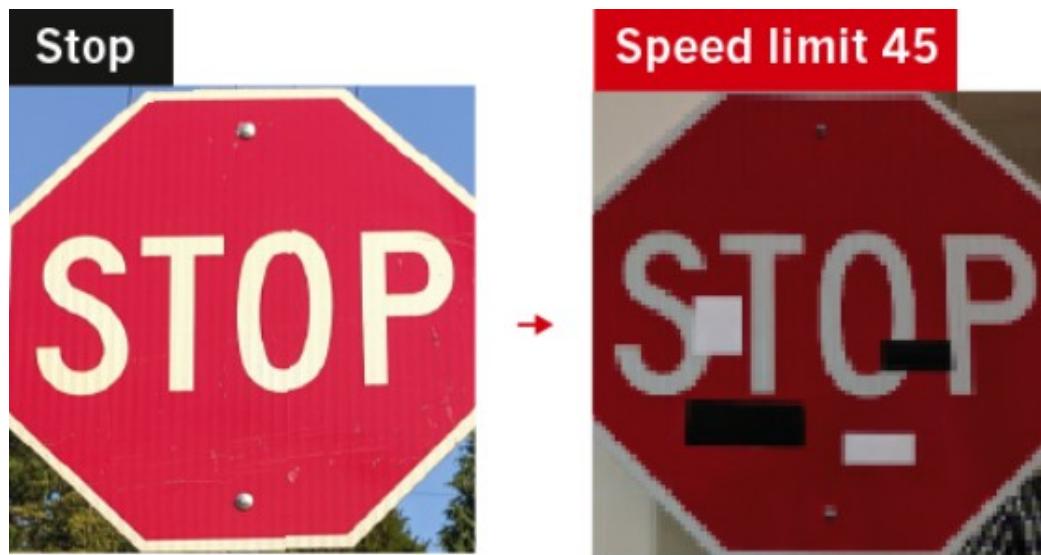


Predicted: **wolf**
True: **husky**

Part II-1 : the deep networks

- **Security**

- Training a network with a biased dataset not only reduce its reliability, but also makes it more vulnerable to attacks



Part II-1 : the deep networks

- **Security**

- It is also possible to ‘manipulate’ the network to modify its output

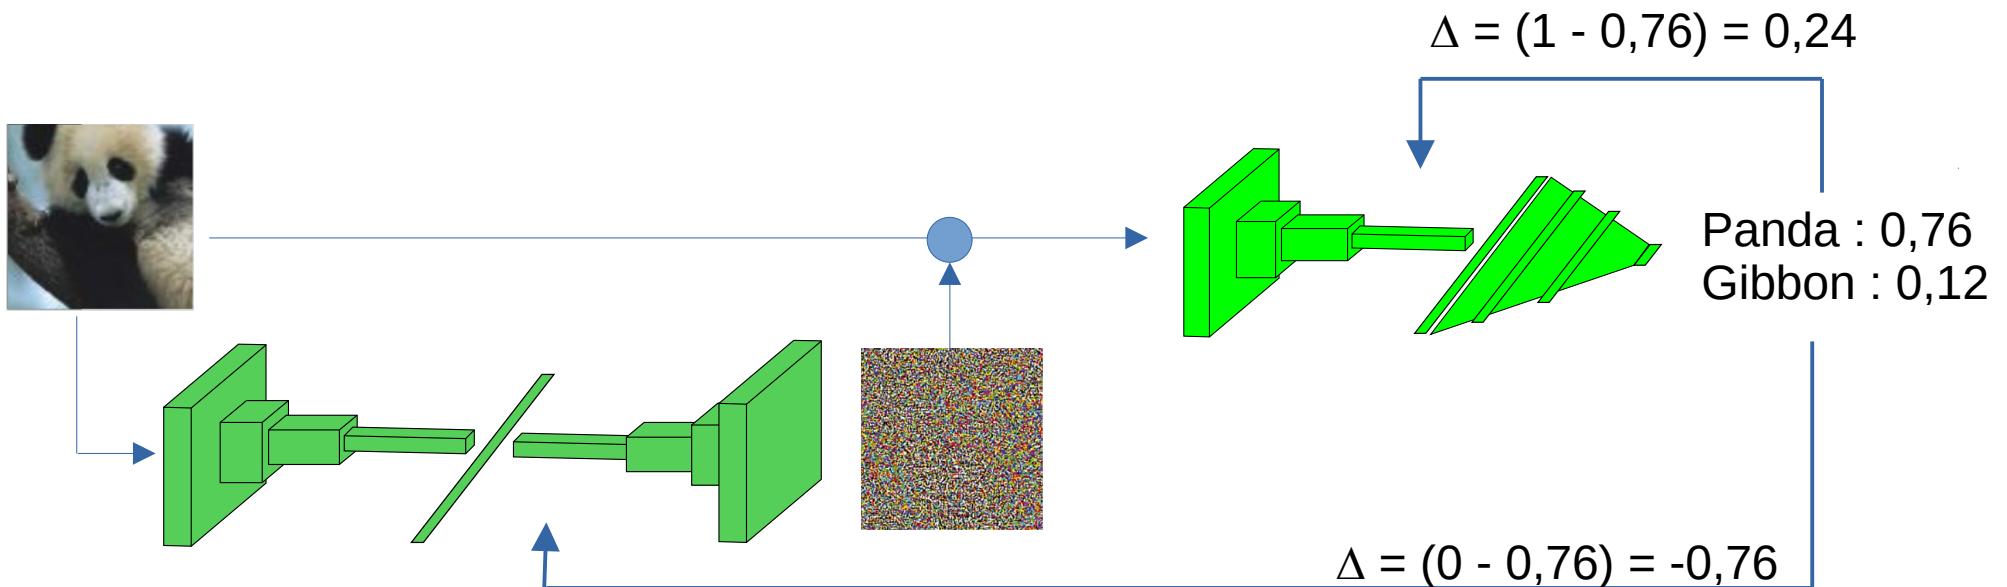


- The added noise disrupted low level features, which also interfere with higher level features, and thus, changed the result.
 - How this variation was created ?

Part II-1 : the deep networks

- **Manipulating a network**

- A ‘normal’ deep network recognize animals on images
- An encoder-decoder network is added to generate a ‘noise image’ that is added to the initial image
- The success of second network is related to the failure of first network



- If the two networks learn simultaneously, they have antagonist objectives !